

Token-Accounting Integrity: Measuring Client-Side Under-Payment in LLM Metering Architectures

(anonymized for review)

Abstract

Usage-based billing is a security boundary in LLM inference: value is streamed to the client before accounting is necessarily finalized, and the charge depends on a multi-category usage record produced at runtime. We study client-side under-payment under an honest-provider threat model, and model this accounting-security boundary along three dimensions — **state synchronization**, **commitment timing**, and **usage authority** — unified by a single integrity property, $V(r) \leq N(r)$, over a request-lifecycle model. *We do not claim new attack primitives*: the enabling mechanisms are individually known, and the contribution is a systematic measurement and architectural analysis of them across representative metering designs, together with an ablation of the enforcement primitives. In controlled experiments the known synchronization race is *created* by concurrency (slope 0.099 per additional concurrent request, $R^2 = 1.0000$), whereas commitment-timing and usage-authority failures remain per-request phenomena: under fixed request volume in the tested architecture their per-request leakage was unchanged across concurrency levels 1–100. For commitment timing, an ablation shows that abort-safe finalization closes accounting-integrity violations while reservation independently provides solvency, and that neither substitutes for the other. For usage authority, computing an accurate server-side recount is insufficient unless that value becomes the billing basis: an architecture that recounts correctly, records the result, and still bills the client-declared number leaks 58.3% of the delivered value. Server-authoritative billing eliminates the observed usage-authority leakage. To test whether these are artifacts of one execution or storage design, we repeat the experiments across three topologies (single worker; a reverse proxy over four workers; a load balancer over two gateway instances) and two accounting backends (a mutable balance row; an append-only ledger with a derived balance): 96 accounting cells match the single-worker control with no mismatches, and 17/17 strongest cases are identical across backends. We also state model-level sufficiency conditions per mechanism, separating what is sufficient in the model from what is implementation-specific. Tokenizer cost is real, grows with input length, and varies by 2.9–3.6 $\times$ across engines; in the tested local SmolLM2-135M CPU configuration the recount accounted for 0.014–0.019% of end-to-end inference time. All reported results regenerate from raw data through fail-closed integrity gates and a clean reproduction pipeline. These results position LLM metering as an accounting-security boundary whose synchronization, lifecycle, and authority choices determine both under-payment exposure and defense cost.

1 Introduction

Because the billed unit — the token — is produced by opaque, streamed, long-running inference, the pipeline that authorizes, meters, and commits a charge is security-relevant. Recent work studies dishonest *providers* inflating that accounting [1, 2, 3], economic attacks that drain a victim’s budget [4, 5], and dishonest *intermediaries* that misroute or forge provenance [10, 11]. All three assume

Table 1: Adversary directions in LLM inference accounting. Prior work covers the first three; this paper studies the fourth.

Adversary	Goal	Representative work	Studied here
Provider	over-charge an honest client	CoIn [1], Invisible Tokens [2], Token Inflation [3]	no
Third party	inflate a <i>victim</i> ’s bill	Denial-of-Wallet [4], LLMjacking [5]	no
Intermediary	misroute / forge provenance	AEX [10], gateway-path provenance [11]	no
Client	under-pay for received inference	—	yes

the client is honest, impersonated, or irrelevant. We study the remaining direction (Table 1): a legitimate but dishonest client that under-pays.

We are deliberately adversarial about our own contribution. Each enabling mechanism is already known in isolation: a deployed LLM gateway refunds pre-consumed quota when a final usage record is missing after a client disconnect [6]; “never trust client input in a security decision” is CWE-807, and a metered-token proxy already mitigates an encoding-induced “silent zero-debit” [7]; concurrent quota-decrement races are a classic web-race primitive with a recent CVE [8, 9]; and practitioner guidance already warns that retries and disconnects must not cause token leakage [12]. Accordingly, **M1 and M2 are systematization and measurement contributions, and B0 is a known baseline**; we make no priority claim.

Contributions.

1. **An accounting model** (§4) connecting inference execution, value delivery, authorization, reservation, debit, reconciliation and refund in a request-lifecycle state machine, with a formal integrity property that admits legitimate refunds.
2. **An architecture taxonomy** of three experimentally distinguishable dimensions — synchronization (B0, known baseline), commitment timing (M1), and usage authority (M2) — with the known/systematized/defensive boundaries stated explicitly.
3. **A reproducible, cross-architecture measurement study**: a testbed and methodology quantifying attack success, leakage, leakage efficiency, detectability, invariant violations, latency and throughput — repeated across three execution topologies and two accounting backends, with integrity gates proven to fail closed and an independent re-implementation of the strongest defense.
4. **Defense sufficiency conditions and their integrity/performance cost**: model-level sufficient conditions per mechanism (§8), an ablation identifying which primitive closes which property, and recount cost measured at three levels (standalone tokenizer, gateway, and a local real-model experiment).

2 Threat Model and Accounting Invariant

Honest provider; dishonest client holding valid credentials. The client controls exactly six things, each exercised by our harness: request payloads including any client-declared usage fields; permitted headers; request concurrency; the connection lifecycle (it may abort at any moment); retries; and observation of its own responses and balance. It *cannot* compromise TLS, access the database or Redis, read the gateway filesystem, compromise the inference backend, forge authenticated upstream metadata, or alter server code. This is a concrete application-layer adversary, not a symbolic

cryptographic one. Every leak we report is therefore attributable to the provider’s own accounting design rather than to a compromise.

For a request r , let $V(r)$ be the value actually delivered, $D(r)$ the committed debits, $F(r)$ refunds, and $N(r) = D(r) - F(r)$ the net committed debit. Under-payment and the integrity property are

$$\text{Leak}(r) = V(r) - N(r), \quad \text{Integrity}(r) : V(r) \leq N(r).$$

The property is an inequality on the *net* debit precisely so that legitimate reconciliation is admitted: a client that consumes less than a reservation must be refunded the difference. A refund is legitimate only when it releases the unconsumed portion, $F(r) \leq \text{Res}(r) - V(r)$; the vulnerable “refund the whole reservation on abort” architecture violates exactly this side condition.

Two distinct properties. Accounting integrity is not the only thing a metering system must provide, and conflating the two obscures the defense analysis. We therefore separate:

Accounting integrity $V(r) \leq N(r)$ — the client is never net-charged less than the value it received.

Solvency a request is served only when the system holds sufficient committed or reserved economic capacity to cover it under the system’s policy.

These are orthogonal. A system can bill exactly for what it delivered while still serving clients who cannot pay (integrity holds, solvency fails), or hold funds and then release them (solvency machinery present, integrity fails). In the M1 dimension, *reservation addresses solvency* and *abort-safe finalization addresses accounting integrity*; §7.2 shows empirically that neither substitutes for the other (Table 6).

3 Background and Related Work

LLM APIs meter multi-category token vectors (input, output, cached, reasoning), stream output incrementally, and settle charges asynchronously. Provider-side auditing work establishes that hidden-token accounting can be inflated and is hard to verify [1, 2, 3]; Denial-of-Wallet and LLM-jacking establish victim-side economic exposure [4, 5]; attestation work establishes intermediary-side provenance [10, 11]. Generic quota races are well understood [8, 9]. What we found no prior study of, *in our targeted corpus*, is an evaluation of client-side under-payment under an honest-provider threat model across representative metering architectures with a common measurement and an ablation of the enforcement primitives. This is a bounded negative from a targeted search; it does not establish non-existence.

4 Accounting State Model

We model a metered inference request as a state machine so that each mechanism can be localized to a specific transition, establishing that the taxonomy is *architectural* rather than a list of unrelated implementation bugs.

4.1 Formal model

For a request r let $U_{\text{true}}(r)$ be the usage actually produced, $U_{\text{bill}}(r)$ the usage billed, and $P(\cdot)$ the (exact-decimal) pricing function. Write $C(r) = P(U_{\text{true}}(r))$ for the true cost, $V(r)$ for the value




Figure 1: Request lifecycle. Value exposure begins at the first delivered token (STREAMING); economic commitment is final only at RECONCILED. The interval between them is the exposure window that M1 exploits. B0 corrupts the AUTHORIZED→RESERVED transition (non-atomic); M2 corrupts the *data* used at ACCOUNTED.

actually delivered to the client, $\text{Res}(r)$ for any reservation, $D(r)$ for committed debits, $F(r)$ for refunds, and

$$N(r) = D(r) - F(r)$$

for the *net committed debit*. Under-payment and the integrity property are

$$\text{Leak}(r) = V(r) - N(r), \quad \text{Integrity}(r) : V(r) \leq N(r).$$

Because legitimate reconciliation exists, integrity is an inequality on the *net* debit; a refund is legitimate only when it releases the unconsumed portion,

$$F(r) \leq \text{Res}(r) - V(r).$$

At trial level the harness independently verifies conservation,

$$\sum_{r \in T} N(r) = B_{\text{initial}} - B_{\text{final}},$$

and fails the build on any divergence.

4.2 Mechanism localization

B0 (synchronization). If two concurrent requests read the same authorization state B and both commit a decrement derived from it, one decrement is lost: $\sum_i N(r_i) < \sum_i V(r_i)$. This requires ≥ 2 concurrent transactions, which is why B0 leaks nothing at concurrency 1.

M1 (commitment timing). If $V(r) > 0$ in STREAMING but the exit path taken (ABORTED) does not reach RECONCILED with $N(r) \geq V(r)$, then $\text{Leak}(r) = V(r) > 0$. The defect is in the transition graph, not in a shared resource, hence it is concurrency-invariant.

M2 (usage authority). If the billing basis is client-controlled, $U_{\text{bill}} = T(U_{\text{true}})$ for an attacker transformation T , then $\text{Leak}(r) = P(U_{\text{true}}) - P(T(U_{\text{true}})) > 0$ whenever $P \circ T < P$. Authority is a per-request data-provenance property, hence also concurrency-invariant.

This predicts the empirical result: B0 leakage is *created* by concurrency, while M1/M2 per-request leakage is invariant to it.

Table 2: Structural distinction between the three mechanisms.

Mechanism	Defective element	Needs concurrency?	Closed by
B0	transition atomicity	yes (≥ 2)	atomic compare-and-decrement
M1	transition graph (abort path)	no	abort-inclusive finalization
M2	data authority at ACCOUNTED	no	server-authoritative recount

5 Taxonomy and Testbed

The three dimensions are experimentally distinguishable, not merely conceptual: **B0** (synchronization) is a non-atomic authorize-then-decrement, is *concurrency-created*, and is labelled a **known baseline** used to validate the measurement rig and to provide a generic reference point; **M1** (commitment timing) is lifecycle-sensitive; **M2** (usage authority) is authority- and pricing-sensitive.

On LLM specificity. We do *not* claim these mechanisms are uniquely possible in LLM systems. Streaming responses, trusting client-supplied values, and racing a counter all occur elsewhere. What is specific to LLM inference is that a single request combines, in one metering boundary: progressive (streamed) value delivery; *runtime-generated* usage that is not known when the request is authorized; multi-category token vectors (input, output, cached, reasoning); tokenizer- and provider-dependent usage semantics; pricing functions defined over those usage categories; and asynchronous reconciliation or settlement. The mechanisms are therefore not claimed as unique security primitives; rather, LLM inference combines the lifecycle, usage-authority and token-vector accounting conditions into one boundary, and it is that combination we systematize and measure. On this basis B0 remains a *generic* baseline — a shared-counter race that behaves identically in any metered API — while M1 and M2 are LLM-relevant accounting dimensions: M1 is only interesting where value is delivered progressively over a long-running request, and M2 only where the billed quantity is a runtime-produced, multi-category vector priced by a category-sensitive function.

The testbed is a FastAPI gateway over PostgreSQL 16 (READ COMMITTED) and Redis 7, with a deterministic mock generator for reproducibility and *real* tokenizers on the recount path, run under Docker Compose on one commodity host (16 logical CPUs, 16 GB). Money is exact decimal. Five M1 architectures and six M2 architectures share common interfaces so that experiments compare *semantics* rather than unrelated code.

6 Experimental Methodology

Category-wise pricing over three tiers; leakage *efficiency* (fraction of true value evaded) is the price-invariant headline metric and absolute dollars are labelled scenario values. Sweeps: B0 concurrency $\{1, 2, 5, 10, 20, 30, 50\} \times 30$ repetitions; M1 architecture $\times$ abort position $\times$ concurrency $\{1, 5, 20, 50, 100\}$; M2 architecture $\times$ manipulation $\times$ concurrency. Request volume is held *constant* across concurrency levels so that volume cannot masquerade as vulnerability.

Integrity gates. Every summarizer independently recomputes each record’s leak from primitives, checks invariant-flag consistency, and verifies ledger conservation $\sum_r N(r) = B_{\text{init}} - B_{\text{final}}$ without trusting the application’s own bookkeeping; a violation fails the build. We validate the gates adversarially: corrupting a leak value, a debit, a served flag, a balance, an invariant flag, or deleting a ledger row must each be detected (12/12), alongside metamorphic properties (10/10). The metamorphic suite also caught one of our own unjustified assumptions — we had asserted

leakage efficiency was price-tier invariant, but that holds only under *uniform* price scaling, not under a structural change to the price ratios — which we corrected rather than suppressed.

Independent verification. The strongest M2 defense is re-implemented in a module sharing no code path with the gateway; verdicts agree on all cross-validated cells (42 M2, 16 M1, 2 B0). The independent checker initially reported three disagreements, which were *definitional*, not defects: it aggregated under-payment only, while the gateway aggregates the signed billing delta. Under flat-total pricing an honest client is systematically *overcharged*, so the two aggregates legitimately differ. We now report both, and use under-payment for the security property.

Statistics. All statistics are produced by a standalone script (scipy/statsmodels) which writes a single output file that the paper cites. Many conditions here are deterministic by construction and have exactly zero variance; for those we report the constant and complete separation rather than a p-value, and reserve nonparametric tests and bootstrap intervals for quantities that genuinely vary.

7 Results

7.1 B0 — synchronization (known baseline)

The non-atomic check-then-decrement leaks only when concurrency ≥ 2 : mean \$-leak is 0.000 at $c=1$, then 0.099, 2.871 and 4.851 at $c=2, 30$ and 50 , growing exactly linearly (slope 0.099 per additional concurrent request, $R^2 = 1.0000$, $p = 1.9 \times 10^{-50}$). The hardened atomic compare-and-decrement leaks \$0 at every level (Table 3, Fig. 5). B0 is not a contribution; it validates the rig and provides the contrast used below.

7.2 M1 — commitment timing

Leakage tracks the value delivered before the missing commit, rising monotonically with abort position and collapsing to zero at completion when the debit finally fires: per request 0.0045 at ~ 1 token delivered, 0.0285 at 25%, 0.0540 at 50%, 0.0960 at 90%, and 0.0000 at 100% (Fig. 2, Table 4). Request-ASR is 1.00 at every abort position before completion. The two vulnerable architectures — “debit only after completion” and “refund the whole reservation on abort” — are economically identical.

M1 comprises two orthogonal security properties. *Property A (solvency):* served value must not exceed the affordable balance. *Property B (accounting finality):* delivered value must not exceed the net committed debit. The ablation separates them (Table 6). With a budget for two requests and six issued, `no_reserve_settle` preserves accounting integrity (leak 0.0000 at every abort position) yet serves all six and drives the balance to -0.32 ; conversely `reserve_refund_on_abort` holds a reservation but refunds it away, failing both. **Reservation and abort-safe finalization are orthogonal controls; neither substitutes for the other.**

Concurrency is not the mechanism. Under fixed request volume in the tested architecture, per-request leakage remained unchanged across concurrency levels 1–100: leak per request measured 0.0528 at $c = 1, 5, 20, 50$ and 100 (identical to four decimal places), and request-ASR was exactly 1.000 (vulnerable) or 0.000 (safe) at every level (Table 5, Fig. 3). The honest-client control arm leaked zero everywhere. We do not claim invariance as a proven property of all systems under all

loads; we report the measured behaviour of this architecture and, in §10, the structural reason the state model predicts it.

7.3 M2 — usage authority

Writing U for the true usage vector, P for the pricing function and T for the attacker’s transformation of the declared usage, a manipulation pays off exactly when $P(T(U)) < P(U)$ — that is, when the pricing basis *reads the field T changes*. Table 11 derives this per manipulation: a total/subtotal mismatch is inert under category pricing (which never reads the declared total) and is the strongest attack under flat-total pricing (0.741), while reclassifying output into a discounted cached category is the reverse. The exploitable set is a property of the billing function, not of attacker ingenuity. Every server-authoritative architecture drives leakage efficiency to 0.000 across all manipulations, and M2 is likewise concurrency-invariant (Table 10).

Usage computation is not usage authority. This is the central M2 result. The decisive case is `client_logged`, which *computes* a correct server-side recount, *records* it, and nevertheless continues to use the client-declared number as the billing basis. Against a 90% output under-report it leaks at efficiency **0.583** — indistinguishable from the `client` architecture that never recounts at all (Table 9). The ablation reproduces the effect independently with a different probe prompt, at 0.595 (Table 14); the two differ only because efficiency depends on the prompt’s token mix, and both show the recount changing nothing. The implication is that

$$\text{correct usage computation} \not\Rightarrow \text{correct billing},$$

and the architectural property actually required is that *the authoritative usage value used for billing must be derived from a trusted source*. Our ablation demonstrates that recounting without authority does not prevent the observed leakage; in the evaluated `client_logged` architecture the recount provides evidence usable for detection or post-hoc reconciliation (level D1, §7.6) but does not prevent the leak. We do not claim this holds for every conceivable post-hoc enforcement design — a system that acted on the recorded discrepancy before settling would behave differently — only that computing the correct number is not by itself sufficient.

Economic sensitivity. Leakage efficiency is exactly invariant under uniform price scaling ($\times 10$) and shifts by at most 0.012 under a price-*structure* change, so the qualitative conclusions are price-independent (Table 12). A side effect of the same asymmetry is that flat-total pricing systematically *overcharges* honest clients (efficiency -0.058).

Concurrency does not amplify these failures. Increased concurrency is not itself a sufficient amplifier of client-side metering evasion. In the tested architecture, B0 exhibited a concurrency-created race, whereas M1 and M2 remained primarily architectural accounting failures. This negative result follows a chain we can state explicitly: the *state model predicts* that a defect in transition *atomicity* depends on contention for shared state, while defects in the transition *graph* (M1) or in *data authority* (M2) are properties of an individual request; that prediction yields the architecture-level *hypothesis* that only B0 should scale with concurrency; the *controlled measurement* holds request volume constant so that volume cannot masquerade as vulnerability; and the *observed distinction* matches. Concretely: B0 is concurrency-sensitive because shared authorization/debit state is raced; M1 is concurrency-independent in the tested design because the violation lies in the abort path of an individual lifecycle; and M2 is concurrency-independent in the tested design because

the violation lies in the authority assigned to per-request usage data. Operationally this means an operator cannot infer safety from low traffic, and cannot detect M1/M2 by watching for load spikes.

7.4 Cross-architecture validation

A measurement study on one deployment invites the objection that its conclusions are artifacts of that deployment. We therefore repeated the experiments across three execution topologies and two accounting backends, changing *only* the deployment or the storage architecture — the attack harness, the architectures under test, and the integrity gates are byte-identical across all runs.

Topologies. (i) *single*: one uvicorn worker, no proxy (the control); (ii) *multi-worker*: nginx over **four** uvicorn workers, i.e. four independent event loops sharing one PostgreSQL and one Redis; (iii) *distributed*: an nginx load balancer over **two separate gateway containers**, two workers each, again sharing one datastore. Every ledger row records the serving process, so we can confirm load actually spread rather than assuming it: four distinct worker processes served M1 requests in the multi-worker run (101/164/79/136 requests), and both containers served traffic in the distributed run.

Across **96** accounting cells compared against the control there were **no mismatches** (Table 7). Differences are either exactly zero or bounded by the price of a single delivered token — the disconnect-detection jitter already characterized in our reproduction audit — and never change a verdict. B0 is also preserved: the hardened posture leaks exactly \$0 at every concurrency in every topology, *including across two separate containers*, because atomicity is enforced by the database rather than by process locality; and the distributed vulnerable posture reproduced the control exactly (\$0.8910 at $c=10$, \$4.8510 at $c=50$).

A multi-worker deployment is not free. Building the topology exposed two real defects that a single-worker study cannot reveal, and we report them as findings rather than silently fixing them. First, the B0 runtime posture was *per-process* state, so a configuration change would have reached only one of four workers; it must live in shared state (we moved it to Redis). Second, four workers racing schema creation at boot crashed with a uniqueness violation on `pg_class`; startup must be serialized (we use a PostgreSQL advisory lock). Neither changes the accounting semantics, but both are prerequisites for the results above to be meaningful.

Accounting backends. To test whether the conclusions depend on *how money is stored*, we implemented a second backend behind the same semantic interface. Backend A keeps a mutable `credits.balance` row and obtains atomicity from a conditional `UPDATE ... WHERE balance >= cost`. Backend B never mutates that row: every movement is an immutable row in an append-only ledger and the balance is *derived* as the sum of those entries, with the atomic reserve becoming a conditional append serialized on an anchor row. Running the strongest B0/M1/M2 cases against both, **17 of 17 were byte-identical** — same leakage, same leakage efficiency, same invariant-violation counts, same reconciliation verdicts (Table 8). This supports the structural claim: the outcomes follow from where synchronization, commitment and authority sit, not from the storage representation.

7.5 Lifecycle failure modes

Beyond the scripted attack we inject nine lifecycle failures (disconnect at $t=0$ and mid-stream, truncated stream, settle-window race, missing / delayed / duplicated usage metadata, and a recount-engine failure). Every invariant violation falls on a vulnerable architecture; the hardened architectures survive all of them; and a recount-engine failure is *fail-closed* — the request is rejected and nothing is served or billed (Table 15).

7.6 Detectability

Detectability is orthogonal to exploitability: a failure can be exploitable and hidden, exploitable but reconcilable, or prevented outright. We classify D0 (no evidence), D1 (post-hoc reconciliation only), D2 (logged discrepancy), D3 (prevented in real time). Under **client** billing the leak is D0 — the ledger faithfully records the client’s lie. Merely *logging* a recount moves the same leak to D1 without preventing it, and reconciling in-request reaches D3 (Table 16). This gives operators who cannot yet re-architect billing a cheap intermediate step.

8 Defense Sufficiency Conditions

The ablation shows which primitive is load-bearing empirically; here we state, for each mechanism, conditions that are *sufficient within the accounting model* to guarantee $V(r) \leq N(r)$. These are model-level deductive arguments, **not** machine-checked proofs and **not** universal theorems about deployed systems; we separate what the model forces from what our implementation happens to supply.

M1 — commitment timing. Sufficient: *(M1-a) terminal-path totality* — every terminal path, including client disconnect, server exception and stream truncation, reaches a state where an accounting event for r is committed and durable; and *(M1-b) refund boundedness* — every refund satisfies $F(r) \leq \text{Res}(r) - V(r)$. If a reservation was taken, (M1-b) gives $N(r) = \text{Res}(r) - F(r) \geq V(r)$; if not, (M1-a) forces the terminal event to commit $D(r) = V(r)$. Solvency is a *separate* condition *(M1-c)*: reserve atomically before execution and refuse the request when the reservation fails. **no_reserve_settle** satisfies (M1-a)/(M1-b) but not (M1-c), and measured exactly that — integrity held (leak 0.0000) while the balance went to -0.32 .

M2 — usage authority. Sufficient: *(M2-a) authority* — the usage value used for billing is derived from state the attacker cannot control, i.e. $U_{\text{bill}} = U_{\text{true}}$; and *(M2-b) pricing monotonicity* — P is non-decreasing in each category. Computing U_{true} is *not* sufficient: an architecture that computes it correctly, records it, and bills $P(U_{\text{decl}})$ satisfies “the server recounts” while violating (M2-a), and leaks $P(U_{\text{true}}) - P(T(U_{\text{true}}))$.

B0 — synchronization. Sufficient: *(B0-a)* the check balance $\geq$ cost and the decrement occur as one operation atomic with respect to competing transactions. Under (B0-a) the authorizations on an account are totally ordered, so no two requests authorize against the same funds. With a single in-flight transaction (B0-a) holds vacuously, which is why B0 leaks nothing at concurrency 1.

Implementation-specific caveats. The model says “committed and durable”; our implementation achieves it with a cancellation-shielded settlement task, and a framework that cancelled that task would satisfy the architecture on paper while violating (M1-a). The **upstream** architecture

satisfies (M2-a) *only because the provider is honest by assumption*. And (B0-a) is a property, not a mechanism: our two backends obtain it by a conditional UPDATE and by an anchor-row lock respectively, which is precisely what §7.4 tests. Full statements, arguments and caveats are in the accompanying `defense_sufficiency` document.

9 Defense Evaluation

Each vulnerable architecture is paired with a server-side primitive: atomic compare-and-decrement (B0), reserve-then-abort-safe-finalization (M1), and server-authoritative recount (M2). Under every defense the invariant-violation rate and leakage efficiency are zero across the full attack catalogue.

What the strongest defense costs. We measure recount at three levels because they answer different questions.

1. **Standalone tokenizer** (Fig. 6): cost is linear in input length (log-log slope 0.999–1.14, $R^2 \geq 0.998$), so it budgets as $\text{tokens} \times \text{a per-token constant}$ ($\approx 0.7 \mu\text{s}/\text{token}$ for tiktoken, $\approx 2.5 \mu\text{s}/\text{token}$ for SentencePiece). At 16,384 tokens the engines span 10.1–41.8ms, a 2.9–3.6 $\times$ implementation effect; at a fixed token count, workload changes cost by up to 3.4 $\times$.
2. **Gateway with a *mock* generator**: recount is nearly free at $\leq 1,024$ tokens ($\geq 0.96\times$ throughput) but costs heavily at 16,384 tokens ($0.05\text{--}0.76\times$ at $c \in \{1, 10\}$). Tokenization time explains this (Spearman $\rho = -0.868$, $p = 1.3 \times 10^{-5}$).
3. **Local real-model experiment**: with real autoregressive generation the recount is **0.014%–0.019%** of end-to-end time across 128-, 1024- and 4096-token contexts, and measured throughput ratios scatter around 1.0 with signs that flip across workloads, i.e. within a $\pm 10\%$ noise band that is itself $\sim 700\times$ larger than the effect.

The apparent expense in (2) is an artifact of the mock generator costing *nothing*: when generation is free, tokenization is the entire request. We report all three rather than the most convenient one, and we do *not* claim recount is “essentially free”: it is a real, predictable cost that is small *relative to real inference* in the tested configuration.

10 Discussion

Three dimensions of one accounting-security boundary behave differently under load, and the difference is predicted by *which* element of the request lifecycle is defective: transition atomicity (B0) is a shared-resource property and was therefore concurrency-created, whereas the transition graph (M1) and data authority (M2) are per-request properties and showed no concurrency dependence in the tested design. The practical consequences are that (i) an operator cannot use traffic level as a proxy for exposure, (ii) integrity and solvency require different controls, and (iii) in the evaluated architectures a correct recount did not prevent leakage unless it was authoritative for billing. The defenses themselves are ordinary; the contribution is knowing precisely which one is load-bearing and what it costs.

“Is this not just CWE-807 plus a race plus known streaming bugs?” This is the objection we consider most likely, and it is partly correct — which is why we state it ourselves. B0 *is* a known generic race and is presented as a baseline, not a contribution. The enabling mechanisms behind M1 and M2 are individually known: trusting client-supplied values is CWE-807, and

Table 3: B0 baseline: credit-decrement race (30 reps/cell).

posture	conc	served	\$-leak	trial ASR	req ASR	detect
hardened	1	1.0	0.0000	0.00	0.000	none
hardened	2	2.0	0.0000	0.00	0.000	none
hardened	5	5.0	0.0000	0.00	0.000	none
hardened	10	10.0	0.0000	0.00	0.000	none
hardened	20	10.0	0.0000	0.00	0.000	none
hardened	30	10.0	0.0000	0.00	0.000	none
hardened	50	10.0	0.0000	0.00	0.000	none
vulnerable	1	1.0	0.0000	0.00	0.000	none
vulnerable	2	2.0	0.0990	1.00	0.500	partial
vulnerable	5	5.0	0.3960	1.00	0.800	partial
vulnerable	10	10.0	0.8910	1.00	0.900	partial
vulnerable	20	20.0	1.8810	1.00	0.950	partial
vulnerable	30	30.0	2.8710	1.00	0.967	partial
vulnerable	50	50.0	4.8510	1.00	0.980	partial

disconnect-related accounting gaps appear in a deployed gateway issue [6], an operations writeup, and practitioner guidance [12]. **We claim no novelty at the primitive level.** What the paper adds is: (i) a unified accounting model in which these otherwise-unrelated failures are the same integrity property violated at different points of one request lifecycle; (ii) a taxonomy that separates synchronization, lifecycle commitment, and usage authority as *experimentally distinguishable* dimensions — the concurrency contrast in §7 is evidence that the separation is real rather than editorial; (iii) an M1 ablation showing that solvency and accounting integrity are orthogonal and require different primitives, which is not derivable from “there is a streaming bug”; (iv) an M2 ablation isolating computation from authority, where an architecture that recounts correctly still leaks 0.583; and (v) a common measurement — economic leakage, leakage efficiency, detectability, invariant violations, and defense overhead — applied uniformly across all three dimensions, which no single prior source provides. We do not assert that practitioners believed a recount was sufficient; we show experimentally that in the evaluated architecture it was not.

11 Threats to Validity

Internal. Single uvicorn worker, so gateway recount numbers are an upper bound; shared host and Docker scheduling; local loopback adds a ~ 45 ms floor per request on the Windows/Docker Desktop configuration, which makes the gateway ratios conservative for small inputs; mock generation everywhere except the real-model experiment; leakage that depends on delivered tokens is sensitive to disconnect timing at high concurrency (the within-cell standard deviation of delivered tokens is exactly zero at $c \leq 20$ and ≈ 0.011 at $c \in \{50, 100\}$).

Table 4: M1 metering-commit timing: leak vs. abort (medium tier, concurrency=1).

architecture	abort%	delivered	leak/req	req ASR	inv.viol	detect
pre_debit	0	1	0.0000	0.00	0.00	n/a
pre_debit	25	17	0.0000	0.00	0.00	n/a
pre_debit	50	34	0.0000	0.00	0.00	n/a
pre_debit	75	52	0.0000	0.00	0.00	n/a
pre_debit	90	62	0.0000	0.00	0.00	n/a
pre_debit	100	69	0.0000	0.00	0.00	n/a
reserve_reconcile	0	1	0.0000	0.00	0.00	n/a
reserve_reconcile	25	17	0.0000	0.00	0.00	n/a
reserve_reconcile	50	34	0.0000	0.00	0.00	n/a
reserve_reconcile	75	52	0.0000	0.00	0.00	n/a
reserve_reconcile	90	62	0.0000	0.00	0.00	n/a
reserve_reconcile	100	69	0.0000	0.00	0.00	n/a
post_completion	0	1	0.0045	1.00	1.00	D0
post_completion	25	17	0.0285	1.00	1.00	D0
post_completion	50	34	0.0540	1.00	1.00	D0
post_completion	75	52	0.0810	1.00	1.00	D0
post_completion	90	62	0.0960	1.00	1.00	D0
post_completion	100	69	0.0000	0.00	0.00	n/a
reserve_refund_on_abort	0	1	0.0045	1.00	1.00	D0
reserve_refund_on_abort	25	17	0.0285	1.00	1.00	D0
reserve_refund_on_abort	50	34	0.0540	1.00	1.00	D0
reserve_refund_on_abort	75	52	0.0810	1.00	1.00	D0
reserve_refund_on_abort	90	62	0.0960	1.00	1.00	D0
reserve_refund_on_abort	100	69	0.0000	0.00	0.00	n/a

Two different high-concurrency questions. It is important to separate *accounting-integrity measurements* from *gateway performance measurements*, because only the latter are affected here.

(a) *Accounting integrity* was measured at $c = 1, 5, 20, 50$ and 100 and **none of it was discarded**: leakage, leakage efficiency, request-ASR and invariant violations are deterministic functions of the ledger, not of throughput, and every trial passed ledger conservation at every concurrency level.

(b) *Gateway throughput ratios* at $c \geq 50$ are excluded from headline performance claims because the benchmark host became the limiting factor. The repeat spread reaches 17% of the median, the no-recount baseline is itself non-monotonic, and single runs produced ratios > 1 — physically impossible for strictly-added work — which is what motivated the switch to median-of-three repeats. These are symptoms of single-worker server and client/event-loop saturation, not of recount cost. Only $c = 1$ and $c = 10$ are treated as performance evidence; the excluded cells remain published in **results/** and summarized here rather than removed.

Construct. Synthetic price tiers; a token-count abstraction of “inference value”; leakage normalized by delivered value; and a detectability scale that separates exploitability from evidence availability.

Table 5: M1: leak PER REQUEST vs concurrency (request volume held constant, attacker arm).

architecture	c=1	c=5	c=20	c=50	c=100
post_completion	0.0528	0.0528	0.0528	0.0528	0.0528
pre_debit	0.0000	0.0000	0.0000	0.0000	0.0000
reserve_reconcile	0.0000	0.0000	0.0000	0.0000	0.0000
reserve_refund_on_abort	0.0528	0.0528	0.0528	0.0528	0.0528

Table 6: M1 architecture matrix. Solvency (served value must not exceed the affordable balance) and accounting integrity (delivered value must not exceed the net committed debit) are ORTHOGONAL: reservation provides the former and abort-safe finalization the latter, and neither substitutes for the other.

architecture	Property A: solvency	Property B: accounting integrity	how the violated property fails
pre_debit	HOLDS	HOLDS	-
reserve_reconcile	HOLDS	HOLDS	-
reserve_refund_on_abort	VIOLATED	VIOLATED	served 6/6 on a 2-request budget
no_reserve_settle	VIOLATED	HOLDS	served 6/6 on a 2-request budget; balance -0.320000
post_completion	VIOLATED	VIOLATED	served 6/6 on a 2-request budget

External. The cross-architecture experiments (§7.4) remove the single-event-loop, single-process and single-storage-design confounds, but they do not remove all of them, and we state the boundary precisely. *Addressed:* four independent event loops; two separate gateway containers behind a load balancer; and two structurally different accounting backends (mutable row vs. append-only derived ledger). *Not addressed:* all instances still run on **one physical host** sharing **one** PostgreSQL and **one** Redis, so multi-host networking, replica lag, cross-region latency and datastore partitioning are untested; both backends are strongly-consistent single-datastore designs, so an **eventually-consistent** or sharded accounting pipeline — which could break the atomicity condition (B0-a) or the durability condition (M1-a) — remains unevaluated. We also test one model family, one cache, and no commercial provider; the real-model experiment bounds a *ratio* in one configuration and does not validate any commercial provider’s billing behaviour.

Formal. The sufficiency conditions of §8 are model-level deductive arguments, not machine-checked proofs. We do not claim they are sufficient in an arbitrary deployed system, nor that no other sufficient condition set exists. Their *necessity* evidence is empirical: removing a condition produced a measured violation in our implementation.

Literature. A targeted search rather than a proof of absence, in a fast-moving area with terminology mismatch risk.

12 Responsible Disclosure and Ethics

This work produced no new vulnerability findings against any third-party, live, or production system; all experiments ran against the local testbed shipped with the artifact, whose harness defaults to `localhost` and contains no third-party endpoints. No credentials, real user data, or production billing APIs were used. The public artifacts we cite as prior evidence (a GitHub issue [6],




Figure 2: M1: leakage vs. abort position (concurrency 1, medium tier). Vulnerable architectures peak just before the commit and collapse to \$0 at completion; safe architectures stay at \$0. Error bars are within-cell standard deviations.

a PR [7], and a CVE [9]) were public before this work and are not our findings. Every attack in the artifact is paired with a defense, and the testbed is released so operators can measure their own gateways.

13 Conclusion

Client-side under-payment in LLM metering is best modelled as one accounting boundary with three placements: synchronization, commitment timing, and usage authority. Localizing each to a state transition predicts — and our measurements confirm — that only the synchronization defect is concurrency-created, that integrity and solvency are closed by different primitives, and that a server-side recount must be authoritative rather than merely computed. In the tested configuration that defense costs a small, predictable fraction of real inference time.

References

- [1] G. Sun et al. CoIn: Counting the Invisible Reasoning Tokens in Commercial Opaque LLM APIs. arXiv:2505.13778, 2025.
- [2] G. Sun et al. Invisible Tokens, Visible Bills. arXiv:2505.18471, 2025.
- [3] S. Hoque et al. Token Inflation: How Dishonest Providers Can Overcharge for LLM Usage. arXiv:2605.30040, 2026.
- [4] OWASP. LLM10: Unbounded Consumption (Denial of Wallet). OWASP Top 10 for LLM Applications, 2025.
- [5] Sysdig Threat Research Team. LLMjacking: stolen cloud credentials used for LLM inference, 2024.
- [6] QuantumNous/new-api issue #5235: streaming requests may refund all pre-consumed quota when final usage is missing after client disconnect, 2026.



Figure 3: M1 negative result: with request volume held constant at 20 requests per cell, per-request leakage and request-ASR are invariant to concurrency from 1 to 100.

- [7] lightninglabs/aperture PR #247: metered draw-down of prepaid L402 tokens; strip **Accept-Encoding** to prevent a silent zero-debit, 2026.
- [8] J. Kettle. The single-packet attack: making remote race-conditions ‘local’. PortSwigger Research, 2023.
- [9] CVE-2026-31873: Tyk API Gateway non-atomic quota check/decrement race, 2026.
- [10] Y. Guan. AEX: Non-Intrusive Multi-Hop Attestation and Provenance for LLM APIs. arXiv:2603.14283, 2026.
- [11] F. Wang and Z. Tian. Evidence-Bound Gateway-Path Provenance for Third-Party LLM Inference. arXiv:2606.22560, 2026.
- [12] Industry metering guidance: gateway as enforcement point vs. meter as accounting point; retries, partial failures and client disconnects must not cause double billing or token leakage, 2026.




Figure 4: M2: leakage efficiency (fraction of delivered value evaded) by architecture $\times$ manipulation, medium tier, concurrency 1. Server-authoritative rows are uniformly zero; the exploitable set differs between category and flat-total pricing.




Figure 5: B0 baseline: mean \$-leak per trial vs. concurrency (30 repetitions per cell). Leakage is *created* by concurrency, in contrast with M1/M2.

Table 7: Cross-topology validation. Accounting outcomes are compared against the single-worker control. 'diff in tokens' expresses an M1 difference in whole delivered tokens; sub-token differences reflect disconnect-detection timing, not accounting divergence. 'distinct procs' confirms load actually spread.

mech	cell	topology	single	topology value	abs diff	diff in tokens	distinct procs	verdict
M1	post_completion @0%	multiworker	0.0045	0.0047	0.0002	0.15	4	WITHIN 1 TOKEN
M1	post_completion @25%	multiworker	0.0285	0.0285	0.0000	0.00	4	IDENTICAL
M1	post_completion @50%	multiworker	0.0540	0.0540	0.0000	0.01	4	WITHIN 1 TOKEN
M1	post_completion @75%	multiworker	0.0810	0.0810	0.0000	0.00	4	IDENTICAL
M1	post_completion @90%	multiworker	0.0960	0.0960	0.0000	0.00	4	IDENTICAL
M1	post_completion @100%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	pre_debit @0%	multiworker	-0.1020	-0.1019	0.0001	0.09	4	WITHIN 1 TOKEN
M1	pre_debit @25%	multiworker	-0.0780	-0.0780	0.0000	0.00	4	IDENTICAL
M1	pre_debit @50%	multiworker	-0.0525	-0.0525	0.0000	0.00	4	IDENTICAL
M1	pre_debit @75%	multiworker	-0.0255	-0.0255	0.0000	0.00	4	IDENTICAL
M1	pre_debit @90%	multiworker	-0.0105	-0.0105	0.0000	0.00	4	IDENTICAL
M1	pre_debit @100%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @0%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @25%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @50%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @75%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @90%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @100%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_refund_on_abort @0%	multiworker	0.0045	0.0045	0.0000	0.03	4	WITHIN 1 TOKEN
M1	reserve_refund_on_abort @25%	multiworker	0.0285	0.0285	0.0000	0.00	4	IDENTICAL
M1	reserve_refund_on_abort @50%	multiworker	0.0540	0.0540	0.0000	0.00	4	IDENTICAL
M1	reserve_refund_on_abort @75%	multiworker	0.0810	0.0810	0.0000	0.00	4	IDENTICAL
M1	reserve_refund_on_abort @90%	multiworker	0.0960	0.0960	0.0000	0.00	4	IDENTICAL
M1	reserve_refund_on_abort @100%	multiworker	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	post_completion @25%	distributed	0.0285	0.0285	0.0000	0.00	4	IDENTICAL
M1	post_completion @50%	distributed	0.0540	0.0540	0.0000	0.01	4	WITHIN 1 TOKEN
M1	post_completion @90%	distributed	0.0960	0.0960	0.0000	0.03	4	WITHIN 1 TOKEN
M1	pre_debit @25%	distributed	-0.0780	-0.0780	0.0000	0.00	4	IDENTICAL
M1	pre_debit @50%	distributed	-0.0525	-0.0525	0.0000	0.00	4	IDENTICAL
M1	pre_debit @90%	distributed	-0.0105	-0.0105	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @25%	distributed	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @50%	distributed	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_reconcile @90%	distributed	0.0000	0.0000	0.0000	0.00	4	IDENTICAL
M1	reserve_refund_on_abort @25%	distributed	0.0285	0.0285	0.0000	0.00	4	IDENTICAL
M1	reserve_refund_on_abort @50%	distributed	0.0540	0.0540	0.0000	0.00	4	IDENTICAL
M1	reserve_refund_on_abort @90%	distributed	0.0960	0.0960	0.0000	0.00	4	IDENTICAL
M2	client/drop_reasoning	multiworker	0.324	0.324	0.0000	-	4	IDENTICAL
M2	client/honest	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	client/inflate_cached	multiworker	0.313	0.313	0.0000	-	4	IDENTICAL
M2	client/rounding_shave	multiworker	0.036	0.036	0.0000	-	4	IDENTICAL
M2	client/total_mismatch	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	client/under_report_output_50	multiworker	0.324	0.324	0.0000	-	4	IDENTICAL
M2	client/under_report_output_90	multiworker	0.583	0.583	0.0000	-	4	IDENTICAL
M2	client_logged/drop_reasoning	multiworker	0.324	0.324	0.0000	-	4	IDENTICAL
M2	client_logged/honest	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	client_logged/inflate_cached	multiworker	0.313	0.313	0.0000	-	4	IDENTICAL
M2	client_logged/rounding_shave	multiworker	0.036	0.036	0.0000	-	4	IDENTICAL
M2	client_logged/total_mismatch	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	client_logged/under_report_output_50	multiworker	0.324	0.324	0.0000	-	4	IDENTICAL
M2	client_logged/under_report_output_90	multiworker	0.583	0.583	0.0000	-	3	IDENTICAL
M2	client_total/drop_reasoning	multiworker	0.266	0.266	0.0000	-	4	IDENTICAL
M2	client_total/honest	multiworker	-0.058	-0.058	0.0000	-	4	IDENTICAL
M2	client_total/inflate_cached	multiworker	-0.058	-0.058	0.0000	-	4	IDENTICAL
M2	client_total/rounding_shave	multiworker	-0.014	-0.014	0.0000	-	4	IDENTICAL
M2	client_total/total_mismatch	multiworker	0.741	0.741	0.0000	-	4	IDENTICAL
M2	client_total/under_report_output_50	multiworker	0.266	0.266	0.0000	-	4	IDENTICAL
M2	client_total/under_report_output_90	multiworker	0.525	0.525	0.0000	-	4	IDENTICAL
M2	hybrid_reconcile/drop_reasoning	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	hybrid_reconcile/honest	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	hybrid_reconcile/inflate_cached	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	hybrid_reconcile/rounding_shave	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	hybrid_reconcile/total_mismatch	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	hybrid_reconcile/under_report_output_50	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	hybrid_reconcile/under_report_output_90	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	server_recount/drop_reasoning	multiworker	0.000	0.000	0.0000	-	3	IDENTICAL
M2	server_recount/honest	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	server_recount/inflate_cached	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	server_recount/rounding_shave	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	server_recount/total_mismatch	multiworker	0.000	0.000	0.0000	-	3	IDENTICAL
M2	server_recount/under_report_output_50	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	server_recount/under_report_output_90	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	upstream/drop_reasoning	multiworker	0.000	0.000	0.0000	-	3	IDENTICAL
M2	upstream/honest	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	upstream/inflate_cached	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	upstream/rounding_shave	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	upstream/total_mismatch	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	upstream/under_report_output_50	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	upstream/under_report_output_90	multiworker	0.000	0.000	0.0000	-	4	IDENTICAL
M2	client/honest	distributed	0.000	0.000	0.0000	-	4	IDENTICAL
M2	client/total_mismatch	distributed	0.000	0.000	0.0000	-	4	IDENTICAL
M2	client/under_report_output_90	distributed	0.583	0.583	0.0000	-	4	IDENTICAL
M2	client_logged/honest	distributed	0.000	0.000	0.0000	-	4	IDENTICAL

Table 8: Cross-backend validation. The same cases run against a mutable balance row and an append-only ledger with a derived balance, behind one semantic interface. Identical outcomes indicate the results do not depend on the balance-storage architecture.

case	metric	mutable	ledger	viol (mut)	viol (led)	verdict
B0/hardened/c10	leak	0.000000	0.000000	0	0	IDENTICAL
B0/hardened/c50	leak	0.000000	0.000000	0	0	IDENTICAL
B0/vulnerable/c10	leak	0.949500	0.949500	9	9	IDENTICAL
B0/vulnerable/c50	leak	5.169500	5.169500	49	49	IDENTICAL
M1/no_reserve_settle/abort50	leak_per_request	0.000000	0.000000	0	0	IDENTICAL
M1/no_reserve_settle/abort90	leak_per_request	0.000000	0.000000	0	0	IDENTICAL
M1/post_completion/abort50	leak_per_request	0.054500	0.054500	10	10	IDENTICAL
M1/post_completion/abort90	leak_per_request	0.095000	0.095000	10	10	IDENTICAL
M1/reserve_reconcile/abort50	leak_per_request	0.000000	0.000000	0	0	IDENTICAL
M1/reserve_reconcile/abort90	leak_per_request	0.000000	0.000000	0	0	IDENTICAL
M1/reserve_refund_on_abort/abort50	leak_per_request	0.054500	0.054500	10	10	IDENTICAL
M1/reserve_refund_on_abort/abort90	leak_per_request	0.095000	0.095000	10	10	IDENTICAL
M2/client/under_report_output_90	efficiency	0.5942492012779552	0.5942492012779552	10	10	IDENTICAL
M2/client_logged/under_report_output_90	efficiency	0.5942492012779552	0.5942492012779552	10	10	IDENTICAL
M2/client_total/total_mismatch	efficiency	0.7412140575079872	0.7412140575079872	10	10	IDENTICAL
M2/hybrid_reconcile/total_mismatch	efficiency	0.0	0.0	0	0	IDENTICAL
M2/server_recount/under_report_output_90	efficiency	0.0	0.0	0	0	IDENTICAL




Figure 6: Standalone tokenizer recount latency (p50, log-log) by engine, workload and input length, over 100 exact-length-verified cells. This is *not* a gateway throughput result.

Table 9: M2 usage-record authority: leakage efficiency by manipulation (medium tier, concurrency=1).

architecture	manipulation	leak eff.	req ASR	inv.viol	detect
client	honest	0.00	0.00	0.00	n/a
client	under_report_output_50	0.32	1.00	1.00	D0
client	under_report_output_90	0.58	1.00	1.00	D0
client	under_report_input_50	0.01	1.00	1.00	D0
client	drop_reasoning	0.32	1.00	1.00	D0
client	inflate_cached	0.31	1.00	1.00	D0
client	total_mismatch	0.00	0.00	0.00	n/a
client	rounding_shave	0.04	1.00	1.00	D0
client_logged	honest	0.00	0.00	0.00	n/a
client_logged	under_report_output_50	0.32	1.00	1.00	D1
client_logged	under_report_output_90	0.58	1.00	1.00	D1
client_logged	under_report_input_50	0.01	1.00	1.00	D1
client_logged	drop_reasoning	0.32	1.00	1.00	D1
client_logged	inflate_cached	0.31	1.00	1.00	D1
client_logged	total_mismatch	0.00	0.00	0.00	n/a
client_logged	rounding_shave	0.04	1.00	1.00	D1
client_total	honest	0.00	0.00	0.00	n/a
client_total	under_report_output_50	0.27	1.00	1.00	D0
client_total	under_report_output_90	0.53	1.00	1.00	D0
client_total	under_report_input_50	0.00	0.00	0.00	n/a
client_total	drop_reasoning	0.27	1.00	1.00	D0
client_total	inflate_cached	0.00	0.00	0.00	n/a
client_total	total_mismatch	0.74	1.00	1.00	D0
client_total	rounding_shave	0.00	0.00	0.00	n/a
upstream	honest	0.00	0.00	0.00	n/a
upstream	under_report_output_50	0.00	0.00	0.00	n/a
upstream	under_report_output_90	0.00	0.00	0.00	n/a
upstream	under_report_input_50	0.00	0.00	0.00	n/a
upstream	drop_reasoning	0.00	0.00	0.00	n/a
upstream	inflate_cached	0.00	0.00	0.00	n/a
upstream	total_mismatch	0.00	0.00	0.00	n/a
upstream	rounding_shave	0.00	0.00	0.00	n/a
server_recount	honest	0.00	0.00	0.00	n/a
server_recount	under_report_output_50	0.00	0.00	0.00	n/a
server_recount	under_report_output_90	0.00	0.00	0.00	n/a
server_recount	under_report_input_50	0.00	0.00	0.00	n/a
server_recount	drop_reasoning	0.00	0.00	0.00	n/a
server_recount	inflate_cached	0.00	0.00	0.00	n/a
server_recount	total_mismatch	0.00	0.00	0.00	n/a
server_recount	rounding_shave	0.00	0.00	0.00	n/a
hybrid_reconcile	honest	0.00	0.00	0.00	n/a
hybrid_reconcile	under_report_output_50	0.00	0.00	0.00	n/a
hybrid_reconcile	under_report_output_90	0.00	0.00	0.00	n/a
hybrid_reconcile	under_report_input_50	0.00	0.00	0.00	n/a
hybrid_reconcile	drop_reasoning	0.00	0.00	0.00	n/a
hybrid_reconcile	inflate_cached	0.00	0.00	0.00	n/a

Table 10: M2: leakage efficiency vs concurrency (attacker arm; volume held constant).

architecture	c=1	c=5	c=20	c=50	c=100
client	0.228	0.228	0.228	0.228	0.228
client_logged	0.228	0.228	0.228	0.228	0.228
client_total	0.257	0.257	0.257	0.257	0.257
hybrid_reconcile	0.000	0.000	0.000	0.000	0.000
server_recount	0.000	0.000	0.000	0.000	0.000
upstream	0.000	0.000	0.000	0.000	0.000

Table 11: M2 as a transformation on the usage vector. Whether a manipulation pays off is determined by whether the pricing basis reads the field T changes.

architecture	pricing basis	manipulation T	vector change T(U)	leak eff.	why it succeeds / fails
client	category	honest	unchanged	+0.000	T is identity (honest)
client	category	under_report_output_50	output:60->30, total:98->68	+0.324	T lowers a category the basis prices
client	category	under_report_output_90	output:60->6, total:98->44	+0.583	T lowers a category the basis prices
client	category	under_report_input_50	input:8->4, total:98->94	+0.014	T lowers a category the basis prices
client	category	drop_reasoning	reasoning:30->0, total:98->68	+0.324	T lowers a category the basis prices
client	category	inflate_cached	output:60->30, cached_input:0->30	+0.313	T lowers a category the basis prices
client	category	total_mismatch	total:98->24	+0.000	T changes only the declared total; category basis never reads it
client	category	rounding_shave	input:8->7, output:60->57, total:98->94	+0.036	T lowers a category the basis prices
client_total	flat-total	honest	unchanged	-0.058	T is identity (honest)
client_total	flat-total	under_report_output_50	output:60->30, total:98->68	+0.266	T lowers the declared total the basis prices
client_total	flat-total	under_report_output_90	output:60->6, total:98->44	+0.525	T lowers the declared total the basis prices
client_total	flat-total	under_report_input_50	input:8->4, total:98->94	-0.014	T's change is priced at (near) the same rate
client_total	flat-total	drop_reasoning	reasoning:30->0, total:98->68	+0.266	T lowers the declared total the basis prices
client_total	flat-total	inflate_cached	output:60->30, cached_input:0->30	-0.058	T changes only subtotals; flat-total basis never reads them
client_total	flat-total	total_mismatch	total:98->24	+0.741	T lowers the declared total the basis prices
client_total	flat-total	rounding_shave	input:8->7, output:60->57, total:98->94	-0.014	T's change is priced at (near) the same rate

Table 12: Economic sensitivity. Dollar columns are SCENARIO values under synthetic price tiers; leakage efficiency and leaked tokens are the price-independent measures. Efficiency is invariant under uniform price scaling (low->medium) but not under a price-structure change (medium->high alters the output/input ratio).

architecture	manipulation	low (\$)	medium (\$)	high (\$)	leaked tokens	leak eff.	inv. under x10 scale?	d(ef) med->high
client	honest	+0.0000	+0.0000	+0.0000	0	+0.000	yes	+0.000
client_total	honest	-0.0160	-0.1600	-1.9200	-113	-0.058	yes	-0.012
client	under_report_output_50	+0.0900	+0.9000	+9.0000	635	+0.324	yes	+0.004
client_total	under_report_output_50	+0.0740	+0.7400	+7.0800	522	+0.266	yes	-0.009
client	under_report_output_90	+0.1620	+1.6200	+16.2000	1142	+0.583	yes	+0.007
client_total	under_report_output_90	+0.1460	+1.4600	+14.2800	1029	+0.525	yes	-0.006
client	under_report_input_50	+0.0040	+0.0400	+0.2400	28	+0.014	yes	-0.006
client_total	under_report_input_50	-0.0040	-0.0400	-0.7200	-28	-0.014	yes	-0.012
client	drop_reasoning	+0.0900	+0.9000	+9.0000	635	+0.324	yes	+0.004
client_total	drop_reasoning	+0.0740	+0.7400	+7.0800	522	+0.266	yes	-0.009
client	inflate_cached	+0.0870	+0.8700	+8.8200	613	+0.313	yes	+0.008
client_total	inflate_cached	-0.0160	-0.1600	-1.9200	-113	-0.058	yes	-0.012
client	total_mismatch	+0.0000	+0.0000	+0.0000	0	+0.000	yes	+0.000
client_total	total_mismatch	+0.2060	+2.0600	+20.2800	1452	+0.741	yes	-0.003
client	rounding_shave	+0.0100	+0.1000	+0.9600	71	+0.036	yes	-0.001
client_total	rounding_shave	-0.0040	-0.0400	-0.7200	-28	-0.014	yes	-0.012

Table 13: M1 ablation: abort-inclusive finalization closes the integrity violation; reservation independently provides the solvency (funds) guarantee. Both are needed for a correct design.

architecture	reservation	abort-finalization	worst leak/req	integrity	served (budget 2)	solvency
pre_debit	yes	n/a (pre-committed)	-0.0090	held	2/6	ok
reserve_reconcile	yes	yes	+0.0000	held	2/6	ok
reserve_refund_on_abort	yes	no	+0.0845	VIOLATED	6/6	ok
no_reserve_settle	no	yes	+0.0000	held	6/6	NEGATIVE
post_completion	no	no	+0.0845	VIOLATED	6/6	ok

Table 14: M2 ablation: performing a recount is not sufficient; it must be used as the billing basis. client_logged computes the recount and still leaks.

architecture	recount performed	recount used for billing	worst leak eff.	integrity	detection
client	no	no	0.595	VIOLATED	D0
client_logged	yes	no	0.595	VIOLATED	D1
client_total	no	no	0.745	VIOLATED	D0
server_recount	yes	yes	0.000	held	n/a
hybrid_reconcile	yes	yes	0.000	held	n/a
upstream	no	no	0.000	held	n/a

Table 15: Lifecycle failure injection. Every invariant violation occurs on a vulnerable architecture; the hardened architectures survive all injected failures, and a recount-engine failure is fail-closed.

id	failure injected	architecture	delivered	net debit	refund	leakage	invariant
F1	client disconnect mid-stream	post_completion	yes	0.0000	0.0000	0.2375	VIOLATED
F1	client disconnect mid-stream	reserve_reconcile	yes	0.2375	0.2175	0.0000	HELD
F2	client disconnect before first token	post_completion	yes	0.0000	0.0000	0.0275	VIOLATED
F2	client disconnect before first token	reserve_reconcile	yes	0.0275	0.4275	0.0000	HELD
F8	partial/truncated stream near completion	post_completion	yes	0.0000	0.0000	0.4325	VIOLATED
F8	partial/truncated stream near completion	reserve_reconcile	yes	0.4325	0.0225	0.0000	HELD
F9	audit read before settle grace period	post_completion	yes	0.0000	0.0000	0.2375	VIOLATED
F9	audit read before settle grace period	reserve_reconcile	yes	0.2375	0.2175	0.0000	HELD
F5	missing/unknown usage metadata	client	yes	0.6725	0.0000	0.0000	HELD
F5	missing/unknown usage metadata	server_recount	yes	0.6725	0.0000	0.0000	HELD
F7	duplicate usage metadata (replayed request_i	client	yes	0.2750	0.0000	0.3975	VIOLATED
F7	duplicate usage metadata (replayed request_i	server_recount	yes	0.6725	0.0000	0.0000	HELD
F6	delayed/out-of-order usage metadata	client	yes	0.4550	0.0000	0.2175	VIOLATED
F6	delayed/out-of-order usage metadata	server_recount	yes	0.6725	0.0000	0.0000	HELD
F4	tokenizer/recount engine unavailable	bench/server_recount	no	0.0000	0.0000	0.0000	HELD

Table 16: Detectability matrix. D0 invisible; D1 only via post-hoc reconciliation; D2 logged discrepancy; D3 prevented or corrected in real time. *B0's final balance alone cannot reveal the loss because the ledger and the balance are both under-decremented consistently with the surviving write.

mech	architecture	req logs	usage ledger	reconciliation	final balance	real-time	level	detection latency
B0	vulnerable (non-atomic)	no	yes	yes	no*	no	D1	until a reconciliation pass
B0	hardened (atomic CAS)	n/a	n/a	n/a	n/a	yes	D3	0 (prevented)
M1	post_completion	partial	no	no	no	no	D0	never (no record is written)
M1	reserve_refund_on_abort	partial	yes	yes	yes	no	D1	until a reconciliation pass
M1	reserve_reconcile	n/a	n/a	n/a	n/a	yes	D3	0 (prevented)
M2	client	no	no	no	no	no	D0	never (ledger records the lie)
M2	client_logged	no	yes	yes	no	no	D1	until a reconciliation pass
M2	server_recount	n/a	n/a	n/a	n/a	yes	D3	0 (prevented)
M2	hybrid_reconcile	yes	yes	yes	n/a	yes	D3	0 (detected and corrected in-request)

Table 17: Architecture safety summary across all attacks and tiers.

mech	architecture	verdict	max req ASR	best detect
M1	pre_debit	safe	0.00	D3
M1	reserve_reconcile	safe	0.00	D3
M1	post_completion	LEAKS	1.00	D0
M1	reserve_refund_on_abort	LEAKS	1.00	D0
M2	client	LEAKS	1.00	D0
M2	client_logged	LEAKS	1.00	D1
M2	client_total	LEAKS	1.00	D0
M2	upstream	safe	0.00	D3
M2	server_recount	safe	0.00	D3
M2	hybrid_reconcile	safe	0.00	D3